

Milestone Report

Self-Improving Agent Harnesses via Trace-Derived RL Environments for Custom-Kernel Engineering

Chandra Suda Aaditya Nalawade

CS 224R: Deep Reinforcement Learning, Stanford | May 23, 2026

1. Experiments so far

Our overarching objective is an ML-research agent that mines its own logs, diagnoses defects, synthesizes verifiable RL environments for those defects, and trains a lightweight LoRA policy that improves future harness behavior. For the milestone we narrowed the verifier domain to GPU kernel engineering (§2) and built kernel-synth (<https://github.com/chandrasuda/kernel-synth>), a trace-derived environment factory. The pipeline shallow-clones a GitHub repo, walks its Python AST to enumerate every `nn.Module` subclass, and selects a diverse, non-obvious subset per repo using either our custom tool-use agent harness or a deterministic AST heuristic. The heuristic favors modules whose forward methods use unusual operations or whose files import known custom-kernel libraries (Triton, FlashAttention, Mamba SSM). Each selected module is then converted into a self-contained RL environment: a frozen reference implementation lifted from the source repo, an auto-inferred input generator, an editable solution scaffold, a benchmarking harness that times eager versus the candidate and checks numerical equivalence, and machine-readable metadata. A small FastAPI viewer browses the catalog. On top of this we wired a Harbor-style ATIF v1.7 rollout layer that uses `torch.compile` as the reward ceiling and exposes Triton-only agent tools inside a workspace sandbox; baseline and `torch.compile` rollouts already validate against the schema.

1.1 Required experiment

We ran the pipeline end-to-end in deterministic-heuristic mode on five PyTorch repos chosen for mechanism diversity:

repo	py files	kept	avg novelty	top picks
state-spaces/mamba	55	8	0.83	Mamba2, Mamba3, MHA
lucidrains/vit-pytorch	76	8	0.92	DSSA, NaViT, LookViT
lucidrains/x-transformers	27	8	0.96	CoPE, DataDependentAlibi, HyperConnection
openai/whisper	14	3	0.21	(down-weighted: vanilla transformer)
lucidrains/audiolm-pytorch	14	8	0.75	Attend, SoundStream, HubertWithKmeans
total	186	35	0.81	~60k LOC scanned

What worked. The heuristic correctly down-weighted vanilla transformer code (Whisper $\Rightarrow$ avg 0.21) and saturated near 1.0 on novel mechanisms (CoPE, DataDependentAlibi, Mamba2/3, DSSA). This is the right inductive bias for our reward, since we don’t want environments on modules where a fused kernel already exists, as the speedup ceiling collapses. What didn’t. The first converter naïvely copied the source file’s full import block, which broke many environments by pulling in unrelated dependencies. Fix: walk the class body for referenced names, keep only the imports actually used, and transitively lift any module-level helper definitions referenced by the class. Trace-derived environments are only useful if the reference module actually runs.

2. Changes to hypothesis / objective

The closed-loop hypothesis is unchanged: agents improve faster when their RL training environments are mined from real trace/code distributions than when they are hand-designed. What narrowed is the verifier (“general ML-research harness defects” $\rightarrow$ “GPU kernel engineering”), because kernels admit a cheap, fully-automatable, dense reward (numerical equivalence + wall-clock speedup vs `torch.compile`), which is exactly what makes long-horizon RL tractable; general ML-research decisions are subjective, slow to verify, and contaminated by API/cost noise, so they are better attacked once the loop is proven on the kernel sub-domain. Preserved: trace-derived envs, GRPO+LoRA training, evaluation vs base / SFT / prompt-only baselines, ablations over env-derivation vs synthetic. New: a concrete environment factory (kernel-synth), an external expert baseline (`torch.compile`) defining the reward ceiling, and Triton as a much smaller action space than the full agentic one (substantially easier to RL).

3. Steps to completion

(1) Harden the rollout layer: stress-test the per-environment benchmark pipeline, tighten the agent-mode tool-use loop, and improve the auto-inferred input generators so more environments run end-to-end out of the box. (2) Generate trajectories at scale: run the latest models and open-source coding LLMs over the 35 envs plus a held-out 5 repos to produce $N \approx 200$ ATIF trajectories with rewards in $\{-0.1\} \cup [-0.2, 1.5]$. (3) GRPO+LoRA training on Qwen3-1.7B or SmolLM3-3B with reward $r = \text{clip}((t_e - t_s)/(t_e - t_c), -0.2, 1.5)$ (t_e, t_s, t_c = eager/solution/compile latency) gated by correctness; evaluate on held-out repos vs base, prompt-only self-reflection, and SFT. (4) Ablations: (a) heuristic vs agent env-selection, (b) trace-derived envs vs a generic Triton-toy-task curriculum (KernelBench-mini style), (c) `torch.compile` ceiling on/off. (5) Final eval: best-of- N solve rate, mean reward, fraction of held-out modules reaching $r \geq 0.5$, and the t_s/t_c distribution.